

I used Claude for understanding papers, breaking apart the large/complex 4M repo we had to interact with and adapt to use with neural data modalities, and brainstorming evaluation and fine-tuning methods.

You can see a lot of these Claude transcripts are centered around adapting the 4M repository - interacting with a lot of files to understand how we can:

1. Add neural data as registered data modalities for 4M
2. How we can enable masking for some neural trials based on image ID
3. How we can create convenient scripts to train 4M of different sizes and different configs on our modal instance

Aside from the 4M repo, I also used Claude to brainstorm eval and fine-tune methods, including:

1. Understanding brainomni architecture deeply and then using those chats to break apart which parts of the model should be frozen or fine-tuned
2. What linear probes for our tokenizers could indicate what

Throughout this process, my root CLAUDE.md file includes instructions that tell Claude to always run decisions by me, and require me to give it a general plan before it starts executing itself, so I keep an understanding of what is happening in the repo.